

CS 474/574 Machine Learning

4. Support Vector Machines (SVMs)

Prof. Dr. Forrest Sheng Bao
Dept. of Computer Science
Iowa State University
Ames, IA, USA

September 21, 2020

Agenda

- ▶ Perceptron algorithm – Its model differs from SVMs a lot. But it shows that the weight vector is a linear combination of some samples – it resembles SVMs in that sense.

Agenda

- ▶ Perceptron algorithm – Its model differs from SVMs a lot. But it shows that the weight vector is a linear combination of some samples – it resembles SVMs in that sense.
- ▶ The intuition of SVMs: separate similar samples of both classes apart as far as possible.

Agenda

- ▶ Perceptron algorithm – Its model differs from SVMs a lot. But it shows that the weight vector is a linear combination of some samples – it resembles SVMs in that sense.
- ▶ The intuition of SVMs: separate similar samples of both classes apart as far as possible.
- ▶ Deriving the primal form of SVMs, and solving it in KKT conditions

Agenda

- ▶ Perceptron algorithm – Its model differs from SVMs a lot. But it shows that the weight vector is a linear combination of some samples – it resembles SVMs in that sense.
- ▶ The intuition of SVMs: separate similar samples of both classes apart as far as possible.
- ▶ Deriving the primal form of SVMs, and solving it in KKT conditions
- ▶ Dual forms of SVMs and the kernel tricks

Agenda

- ▶ Perceptron algorithm – Its model differs from SVMs a lot. But it shows that the weight vector is a linear combination of some samples – it resembles SVMs in that sense.
- ▶ The intuition of SVMs: separate similar samples of both classes apart as far as possible.
- ▶ Deriving the primal form of SVMs, and solving it in KKT conditions
- ▶ Dual forms of SVMs and the kernel tricks
- ▶ Soft-margin SVMs

All samples are equal. But some samplers are equaler.

- ▶ Let's first see a demo of a linear classifier for linearly separable cases. Pay attention to the prediction outcome.

All samples are equal. But some samplers are equaler.

- ▶ Let's first see a demo of a linear classifier for linearly separable cases. Pay attention to the prediction outcome.
- ▶ Think about the error-based loss function for a classifier: $\sum_i (\hat{y} - y)^2$ where y is the ground truth label and \hat{y} is the prediction.

All samples are equal. But some samplers are equaler.

- ▶ Let's first see a demo of a linear classifier for linearly separable cases. Pay attention to the prediction outcome.
- ▶ Think about the error-based loss function for a classifier: $\sum_i (\hat{y} - y)^2$ where y is the ground truth label and \hat{y} is the prediction.
- ▶ If $y = +1$ and $\hat{y} = +1.5$, should the error be 0.25 or 0 (because properly classified)?

The perceptron algorithm

- ▶ Recall earlier that a sample (\mathbf{x}_i, y_i) is correctly classified if $\mathbf{w}^T \mathbf{x}_i y_i > 0$.

The perceptron algorithm

- ▶ Recall earlier that a sample (\mathbf{x}_i, y_i) is correctly classified if $\mathbf{w}^T \mathbf{x}_i y_i > 0$.
- ▶ Let's define a new cost function to be minimized: $J(\mathbf{w}) = \sum_{x_i \in \mathcal{M}} -\mathbf{w}^T \mathbf{x}_i y_i$ where \mathcal{M} is the set of all samples misclassified ($\mathbf{w}^T \mathbf{x}_i y_i < 0$).

The perceptron algorithm

- ▶ Recall earlier that a sample (\mathbf{x}_i, y_i) is correctly classified if $\mathbf{w}^T \mathbf{x}_i y_i > 0$.
- ▶ Let's define a new cost function to be minimized: $J(\mathbf{w}) = \sum_{x_i \in \mathcal{M}} -\mathbf{w}^T \mathbf{x}_i y_i$ where \mathcal{M} is the set of all samples misclassified ($\mathbf{w}^T \mathbf{x}_i y_i < 0$).
- ▶ Then, $\nabla J(\mathbf{w}) = \sum_{\mathbf{x}_i \in \mathcal{M}} -\mathbf{x}_i y_i$ (because \mathbf{w} is the coefficients.)

The perceptron algorithm

- ▶ Recall earlier that a sample (\mathbf{x}_i, y_i) is correctly classified if $\mathbf{w}^T \mathbf{x}_i y_i > 0$.
- ▶ Let's define a new cost function to be minimized: $J(\mathbf{w}) = \sum_{x_i \in \mathcal{M}} -\mathbf{w}^T \mathbf{x}_i y_i$ where \mathcal{M} is the set of all samples misclassified ($\mathbf{w}^T \mathbf{x}_i y_i < 0$).
- ▶ Then, $\nabla J(\mathbf{w}) = \sum_{\mathbf{x}_i \in \mathcal{M}} -\mathbf{x}_i y_i$ (because \mathbf{w} is the coefficients.)
- ▶ Only those misclassified matter!

The perceptron algorithm

- ▶ Recall earlier that a sample (\mathbf{x}_i, y_i) is correctly classified if $\mathbf{w}^T \mathbf{x}_i y_i > 0$.
- ▶ Let's define a new cost function to be minimized: $J(\mathbf{w}) = \sum_{x_i \in \mathcal{M}} -\mathbf{w}^T \mathbf{x}_i y_i$ where \mathcal{M} is the set of all samples misclassified ($\mathbf{w}^T \mathbf{x}_i y_i < 0$).
- ▶ Then, $\nabla J(\mathbf{w}) = \sum_{\mathbf{x}_i \in \mathcal{M}} -\mathbf{x}_i y_i$ (because \mathbf{w} is the coefficients.)
- ▶ Only those misclassified matter!
- ▶ Batch perceptron algorithm: In each batch, compute $\nabla J(\mathbf{w})$ for all samples misclassified using the same current \mathbf{w} and then update.

Single-sample perceptron algorithm

- ▶ Another common type of perceptron algorithm is called single-sample perceptron algorithm.

Single-sample perceptron algorithm

- ▶ Another common type of perceptron algorithm is called single-sample perceptron algorithm.
- ▶ Update \mathbf{w} whenever a sample is misclassified.

Single-sample perceptron algorithm

- ▶ Another common type of perceptron algorithm is called single-sample perceptron algorithm.
- ▶ Update \mathbf{w} whenever a sample is misclassified.
 1. Initially, \mathbf{w} has arbitrary values. $k = 1$.

Single-sample perceptron algorithm

- ▶ Another common type of perceptron algorithm is called single-sample perceptron algorithm.
- ▶ Update \mathbf{w} whenever a sample is misclassified.
 1. Initially, \mathbf{w} has arbitrary values. $k = 1$.
 2. In the k -th iteration, use sample \mathbf{x}_j such that $j = k \bmod n$ to update the \mathbf{w} by:

$$\mathbf{W}_{k+1} = \begin{cases} \mathbf{W}_k + \rho \mathbf{X}_j y_j & , \text{ if } \mathbf{W}_j^T \mathbf{X}_j y_j \leq 0, \text{ (wrong prediction)} \\ \mathbf{W}_k & , \text{ if } \mathbf{W}_j^T \mathbf{X}_j y_j > 0 \text{ (correct classification)} \end{cases}$$

where ρ is a constant called **learning rate**.

Single-sample perceptron algorithm

- ▶ Another common type of perceptron algorithm is called single-sample perceptron algorithm.
- ▶ Update \mathbf{w} whenever a sample is misclassified.
 1. Initially, \mathbf{w} has arbitrary values. $k = 1$.
 2. In the k -th iteration, use sample \mathbf{x}_j such that $j = k \bmod n$ to update the \mathbf{w} by:

$$\mathbf{W}_{k+1} = \begin{cases} \mathbf{W}_k + \rho \mathbf{X}_j y_j & , \text{ if } \mathbf{W}_j^T \mathbf{X}_j y_j \leq 0, \text{ (wrong prediction)} \\ \mathbf{W}_k & , \text{ if } \mathbf{W}_j^T \mathbf{X}_j y_j > 0 \text{ (correct classification)} \end{cases}$$

where ρ is a constant called **learning rate**.

3. The algorithm terminates when all samples are classified correctly.

Single-sample perceptron algorithm

- ▶ Another common type of perceptron algorithm is called single-sample perceptron algorithm.
- ▶ Update \mathbf{w} whenever a sample is misclassified.
 1. Initially, \mathbf{w} has arbitrary values. $k = 1$.
 2. In the k -th iteration, use sample \mathbf{x}_j such that $j = k \bmod n$ to update the \mathbf{w} by:

$$\mathbf{W}_{k+1} = \begin{cases} \mathbf{W}_k + \rho \mathbf{X}_j y_j & , \text{ if } \mathbf{W}_j^T \mathbf{X}_j y_j \leq 0, \text{ (wrong prediction)} \\ \mathbf{W}_k & , \text{ if } \mathbf{W}_j^T \mathbf{X}_j y_j > 0 \text{ (correct classification)} \end{cases}$$

where ρ is a constant called **learning rate**.

3. The algorithm terminates when all samples are classified correctly.
- ▶ Note that \mathbf{x}_k is not necessarily the k -th training sample due to the loop.

Now let's begin the SVM journey.

An example of single-sample perceptron algorithm

- ▶ Feature vectors and labels:

An example of single-sample perceptron algorithm

- ▶ Feature vectors and labels:

- ▶ $\mathbf{x}'_1 = (0, 0)^T, y_1 = 1$

An example of single-sample perceptron algorithm

- ▶ Feature vectors and labels:

- ▶ $\mathbf{x}'_1 = (0, 0)^T, y_1 = 1$

- ▶ $\mathbf{x}'_2 = (0, 1)^T, y_2 = 1$

An example of single-sample perceptron algorithm

- ▶ Feature vectors and labels:

- ▶ $\mathbf{x}'_1 = (0, 0)^T, y_1 = 1$
- ▶ $\mathbf{x}'_2 = (0, 1)^T, y_2 = 1$
- ▶ $\mathbf{x}'_3 = (1, 0)^T, y_3 = -1$

An example of single-sample perceptron algorithm

- ▶ Feature vectors and labels:

- ▶ $\mathbf{x}'_1 = (0, 0)^T, y_1 = 1$
- ▶ $\mathbf{x}'_2 = (0, 1)^T, y_2 = 1$
- ▶ $\mathbf{x}'_3 = (1, 0)^T, y_3 = -1$
- ▶ $\mathbf{x}'_4 = (1, 1)^T, y_4 = -1$

An example of single-sample perceptron algorithm

- ▶ Feature vectors and labels:
 - ▶ $\mathbf{x}'_1 = (0, 0)^T, y_1 = 1$
 - ▶ $\mathbf{x}'_2 = (0, 1)^T, y_2 = 1$
 - ▶ $\mathbf{x}'_3 = (1, 0)^T, y_3 = -1$
 - ▶ $\mathbf{x}'_4 = (1, 1)^T, y_4 = -1$
- ▶ First, let's augment them and multiply with the labels:

An example of single-sample perceptron algorithm

- ▶ Feature vectors and labels:
 - ▶ $\mathbf{x}'_1 = (0, 0)^T, y_1 = 1$
 - ▶ $\mathbf{x}'_2 = (0, 1)^T, y_2 = 1$
 - ▶ $\mathbf{x}'_3 = (1, 0)^T, y_3 = -1$
 - ▶ $\mathbf{x}'_4 = (1, 1)^T, y_4 = -1$
- ▶ First, let's augment them and multiply with the labels:
 - ▶ $\mathbf{x}_1 y_1 = (0, 0, 1)^T,$

An example of single-sample perceptron algorithm

- ▶ Feature vectors and labels:
 - ▶ $\mathbf{x}'_1 = (0, 0)^T, y_1 = 1$
 - ▶ $\mathbf{x}'_2 = (0, 1)^T, y_2 = 1$
 - ▶ $\mathbf{x}'_3 = (1, 0)^T, y_3 = -1$
 - ▶ $\mathbf{x}'_4 = (1, 1)^T, y_4 = -1$
- ▶ First, let's augment them and multiply with the labels:
 - ▶ $\mathbf{x}_1 y_1 = (0, 0, 1)^T,$
 - ▶ $\mathbf{x}_2 y_2 = (0, 1, 1)^T,$

An example of single-sample perceptron algorithm

- ▶ Feature vectors and labels:
 - ▶ $\mathbf{x}'_1 = (0, 0)^T, y_1 = 1$
 - ▶ $\mathbf{x}'_2 = (0, 1)^T, y_2 = 1$
 - ▶ $\mathbf{x}'_3 = (1, 0)^T, y_3 = -1$
 - ▶ $\mathbf{x}'_4 = (1, 1)^T, y_4 = -1$
- ▶ First, let's augment them and multiply with the labels:
 - ▶ $\mathbf{x}_1 y_1 = (0, 0, 1)^T,$
 - ▶ $\mathbf{x}_2 y_2 = (0, 1, 1)^T,$
 - ▶ $\mathbf{x}_3 y_3 = (-1, 0, -1)^T$

An example of single-sample perceptron algorithm

- ▶ Feature vectors and labels:
 - ▶ $\mathbf{x}'_1 = (0, 0)^T, y_1 = 1$
 - ▶ $\mathbf{x}'_2 = (0, 1)^T, y_2 = 1$
 - ▶ $\mathbf{x}'_3 = (1, 0)^T, y_3 = -1$
 - ▶ $\mathbf{x}'_4 = (1, 1)^T, y_4 = -1$
- ▶ First, let's augment them and multiply with the labels:
 - ▶ $\mathbf{x}_1 y_1 = (0, 0, 1)^T,$
 - ▶ $\mathbf{x}_2 y_2 = (0, 1, 1)^T,$
 - ▶ $\mathbf{x}_3 y_3 = (-1, 0, -1)^T$
 - ▶ $\mathbf{x}_4 y_4 = (-1, -1, -1)^T$

An example of single-sample perceptron algorithm

► Feature vectors and labels:

- $\mathbf{x}'_1 = (0, 0)^T$, $y_1 = 1$
- $\mathbf{x}'_2 = (0, 1)^T$, $y_2 = 1$
- $\mathbf{x}'_3 = (1, 0)^T$, $y_3 = -1$
- $\mathbf{x}'_4 = (1, 1)^T$, $y_4 = -1$

► First, let's augment them and multiply with the labels:

- $\mathbf{x}_1 y_1 = (0, 0, 1)^T$,
- $\mathbf{x}_2 y_2 = (0, 1, 1)^T$,
- $\mathbf{x}_3 y_3 = (-1, 0, -1)^T$
- $\mathbf{x}_4 y_4 = (-1, -1, -1)^T$

0. Begin our iteration. Let $\mathbf{w}_1 = (0, 0, 0)^T$ and $\rho = 1$.

An example of single-sample perceptron algorithm

► Feature vectors and labels:

- $\mathbf{x}'_1 = (0, 0)^T$, $y_1 = 1$
- $\mathbf{x}'_2 = (0, 1)^T$, $y_2 = 1$
- $\mathbf{x}'_3 = (1, 0)^T$, $y_3 = -1$
- $\mathbf{x}'_4 = (1, 1)^T$, $y_4 = -1$

► First, let's augment them and multiply with the labels:

- $\mathbf{x}_1 y_1 = (0, 0, 1)^T$,
- $\mathbf{x}_2 y_2 = (0, 1, 1)^T$,
- $\mathbf{x}_3 y_3 = (-1, 0, -1)^T$
- $\mathbf{x}_4 y_4 = (-1, -1, -1)^T$

0. Begin our iteration. Let

$\mathbf{w}_1 = (0, 0, 0)^T$ and $\rho = 1$.

1. $\mathbf{W}_1^T \cdot \mathbf{x}_1 y_1 = \begin{pmatrix} 0 & 0 & 0 \end{pmatrix} \begin{pmatrix} 0 \\ 0 \\ 1 \end{pmatrix} = 0 \leq 0.$

Need to update \mathbf{W} : $\mathbf{W}_2 =$

$$\mathbf{W}_1 + \rho \cdot \mathbf{x}_1 y_1 = \begin{pmatrix} 0 \\ 0 \\ 0 \end{pmatrix} + \begin{pmatrix} 0 \\ 0 \\ 1 \end{pmatrix} = \begin{pmatrix} 0 \\ 0 \\ 1 \end{pmatrix}$$

An example of single-sample perceptron algorithm

► Feature vectors and labels:

- $\mathbf{x}'_1 = (0, 0)^T$, $y_1 = 1$
- $\mathbf{x}'_2 = (0, 1)^T$, $y_2 = 1$
- $\mathbf{x}'_3 = (1, 0)^T$, $y_3 = -1$
- $\mathbf{x}'_4 = (1, 1)^T$, $y_4 = -1$

► First, let's augment them and multiply with the labels:

- $\mathbf{x}_1 y_1 = (0, 0, 1)^T$,
- $\mathbf{x}_2 y_2 = (0, 1, 1)^T$,
- $\mathbf{x}_3 y_3 = (-1, 0, -1)^T$
- $\mathbf{x}_4 y_4 = (-1, -1, -1)^T$

2. $\mathbf{W}_2^T \cdot \mathbf{x}_2 y_2 = \begin{pmatrix} 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} 0 \\ 1 \\ 1 \end{pmatrix} = 1 > 0$. No updated need. But since \mathbf{w} so far does not classify all samples correctly, we need to keep going. Just let $\mathbf{w}_3 = \mathbf{w}_2$.

0. Begin our iteration. Let

$\mathbf{w}_1 = (0, 0, 0)^T$ and $\rho = 1$.

1. $\mathbf{W}_1^T \cdot \mathbf{x}_1 y_1 = \begin{pmatrix} 0 & 0 & 0 \end{pmatrix} \begin{pmatrix} 0 \\ 0 \\ 1 \end{pmatrix} = 0 \leq 0$.

Need to update \mathbf{W} : $\mathbf{W}_2 =$

$$\mathbf{W}_1 + \rho \cdot \mathbf{x}_1 y_1 = \begin{pmatrix} 0 \\ 0 \\ 0 \end{pmatrix} + \begin{pmatrix} 0 \\ 0 \\ 1 \end{pmatrix} = \begin{pmatrix} 0 \\ 0 \\ 1 \end{pmatrix}$$

An example of perceptron algorithm (cond.)

Continue in perceptron.ipynb

14. In the end, we have $\mathbf{W}_{14} = \begin{pmatrix} -3 \\ 0 \\ 2 \end{pmatrix}$,

let's verify how well it works

An example of perceptron algorithm (cond.)

Continue in perceptron.ipynb

14. In the end, we have $\mathbf{w}_{14} = \begin{pmatrix} -3 \\ 0 \\ 2 \end{pmatrix}$,

let's verify how well it works



$$\begin{cases} \mathbf{w}_{14} \cdot \mathbf{x}_1 y_1 &= 1 > 0 \\ \mathbf{w}_{14} \cdot \mathbf{x}_2 y_2 &= 1 > 0 \\ \mathbf{w}_{14} \cdot \mathbf{x}_3 y_3 &= 1 > 0 \\ \mathbf{w}_{14} \cdot \mathbf{x}_4 y_4 &= 1 > 0 \end{cases}$$

An example of perceptron algorithm (cond.)

Continue in perceptron.ipynb

14. In the end, we have $\mathbf{w}_{14} = \begin{pmatrix} -3 \\ 0 \\ 2 \end{pmatrix}$,

let's verify how well it works

► Mission accomplished!



$$\begin{cases} \mathbf{w}_{14} \cdot \mathbf{x}_1 y_1 &= 1 > 0 \\ \mathbf{w}_{14} \cdot \mathbf{x}_2 y_2 &= 1 > 0 \\ \mathbf{w}_{14} \cdot \mathbf{x}_3 y_3 &= 1 > 0 \\ \mathbf{w}_{14} \cdot \mathbf{x}_4 y_4 &= 1 > 0 \end{cases}$$

An example of perceptron algorithm (cond.)

Continue in perceptron.ipynb

14. In the end, we have $\mathbf{W}_{14} = \begin{pmatrix} -3 \\ 0 \\ 2 \end{pmatrix}$,

let's verify how well it works



$$\begin{cases} \mathbf{w}_{14} \cdot \mathbf{x}_1 y_1 &= 1 > 0 \\ \mathbf{w}_{14} \cdot \mathbf{x}_2 y_2 &= 1 > 0 \\ \mathbf{w}_{14} \cdot \mathbf{x}_3 y_3 &= 1 > 0 \\ \mathbf{w}_{14} \cdot \mathbf{x}_4 y_4 &= 1 > 0 \end{cases}$$

- ▶ Mission accomplished!
- ▶ Note that the perceptron algorithm will not *converge* unless the data is linearly separable.

An example of perceptron algorithm (cond.)

Continue in perceptron.ipynb

14. In the end, we have $\mathbf{W}_{14} = \begin{pmatrix} -3 \\ 0 \\ 2 \end{pmatrix}$,

let's verify how well it works



$$\begin{cases} \mathbf{w}_{14} \cdot \mathbf{x}_1 y_1 &= 1 > 0 \\ \mathbf{w}_{14} \cdot \mathbf{x}_2 y_2 &= 1 > 0 \\ \mathbf{w}_{14} \cdot \mathbf{x}_3 y_3 &= 1 > 0 \\ \mathbf{w}_{14} \cdot \mathbf{x}_4 y_4 &= 1 > 0 \end{cases}$$

- ▶ Mission accomplished!
- ▶ Note that the perceptron algorithm will not *converge* unless the data is linearly separable.
- ▶ What is \mathbf{w} exactly? A linear composition of all training samples!

An example of perceptron algorithm (cond.)

Continue in perceptron.ipynb

14. In the end, we have $\mathbf{w}_{14} = \begin{pmatrix} -3 \\ 0 \\ 2 \end{pmatrix}$,

let's verify how well it works



$$\begin{cases} \mathbf{w}_{14} \cdot \mathbf{x}_1 y_1 &= 1 > 0 \\ \mathbf{w}_{14} \cdot \mathbf{x}_2 y_2 &= 1 > 0 \\ \mathbf{w}_{14} \cdot \mathbf{x}_3 y_3 &= 1 > 0 \\ \mathbf{w}_{14} \cdot \mathbf{x}_4 y_4 &= 1 > 0 \end{cases}$$

- ▶ Mission accomplished!
- ▶ Note that the perceptron algorithm will not *converge* unless the data is linearly separable.
- ▶ What is \mathbf{w} exactly? A linear composition of all training samples!
- ▶ Do all samples contribute to \mathbf{w} ? Not really!

Getting ready for SVMs

- ▶ Earlier our discussion used the augmented definition of linear binary classifier: the feature vector $\mathbf{x} = (x_1, \dots, x_n, 1)^T$ and the weight vector $\mathbf{w} = (w_1, \dots, w_n, w_b)^T$. The hyperplane is an equation $\mathbf{w}^T \mathbf{x} = 0$. If $\mathbf{w}^T \mathbf{x} > 0$, then the sample belongs to one class. If $\mathbf{w}^T \mathbf{x} < 0$, the other class.

Getting ready for SVMs

- ▶ Earlier our discussion used the augmented definition of linear binary classifier: the feature vector $\mathbf{x} = (x_1, \dots, x_n, 1)^T$ and the weight vector $\mathbf{w} = (w_1, \dots, w_n, w_b)^T$. The hyperplane is an equation $\mathbf{w}^T \mathbf{x} = 0$. If $\mathbf{w}^T \mathbf{x} > 0$, then the sample belongs to one class. If $\mathbf{w}^T \mathbf{x} < 0$, the other class.
- ▶ Let's go back to the un-augmented version. Let $\mathbf{x} = [x_1, x_2, \dots, x_n]^T$ and $\mathbf{w} = [w_1, w_2, \dots, w_n]^T$. If $\mathbf{w}^T \mathbf{x} + w_b > 0$ then $\mathbf{x} \in C_1$. If $\mathbf{w}^T \mathbf{x} + w_b < 0$ then $\mathbf{x} \in C_2$. The equation $\mathbf{w}^T \mathbf{x} + w_b = 0$ is the hyperplane, where \mathbf{w} only determines the direction of the hyperplane. To build a classifier is to search for the values for w_1, \dots, w_n and w_b , the bias/threshold.


Getting ready for SVMs

- ▶ Earlier our discussion used the augmented definition of linear binary classifier: the feature vector $\mathbf{x} = (x_1, \dots, x_n, 1)^T$ and the weight vector $\mathbf{w} = (w_1, \dots, w_n, w_b)^T$. The hyperplane is an equation $\mathbf{w}^T \mathbf{x} = 0$. If $\mathbf{w}^T \mathbf{x} > 0$, then the sample belongs to one class. If $\mathbf{w}^T \mathbf{x} < 0$, the other class.
- ▶ Let's go back to the un-augmented version. Let $\mathbf{x} = [x_1, x_2, \dots, x_n]^T$ and $\mathbf{w} = [w_1, w_2, \dots, w_n]^T$. If $\mathbf{w}^T \mathbf{x} + w_b > 0$ then $\mathbf{x} \in C_1$. If $\mathbf{w}^T \mathbf{x} + w_b < 0$ then $\mathbf{x} \in C_2$. The equation $\mathbf{w}^T \mathbf{x} + w_b = 0$ is the hyperplane, where \mathbf{w} only determines the direction of the hyperplane. To build a classifier is to search for the values for w_1, \dots, w_n and w_b , the bias/threshold.
- ▶ For convenience, we denote $g(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$.

Getting ready for SVMs


- ▶ Earlier our discussion used the augmented definition of linear binary classifier: the feature vector $\mathbf{x} = (x_1, \dots, x_n, 1)^T$ and the weight vector $\mathbf{w} = (w_1, \dots, w_n, w_b)^T$. The hyperplane is an equation $\mathbf{w}^T \mathbf{x} = 0$. If $\mathbf{w}^T \mathbf{x} > 0$, then the sample belongs to one class. If $\mathbf{w}^T \mathbf{x} < 0$, the other class.
- ▶ Let's go back to the un-augmented version. Let $\mathbf{x} = [x_1, x_2, \dots, x_n]^T$ and $\mathbf{w} = [w_1, w_2, \dots, w_n]^T$. If $\mathbf{w}^T \mathbf{x} + w_b > 0$ then $\mathbf{x} \in C_1$. If $\mathbf{w}^T \mathbf{x} + w_b < 0$ then $\mathbf{x} \in C_2$. The equation $\mathbf{w}^T \mathbf{x} + w_b = 0$ is the hyperplane, where \mathbf{w} only determines the direction of the hyperplane. To build a classifier is to search for the values for w_1, \dots, w_n and w_b , the bias/threshold.
- ▶ For convenience, we denote $g(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$.
- ▶ We have proved that \mathbf{w} , augmented or not, is perpendicular to the hyperlane.

What is the distance from a sample \mathbf{z} to the hyperplane?




1. Let the point on the hyperplane closest to \mathbf{z} be \mathbf{x} .
Define $\mathbf{y} = \mathbf{x} - \mathbf{z}$.

What is the distance from a sample \mathbf{z} to the hyperplane?




1. Let the point on the hyperplane closest to \mathbf{z} be \mathbf{x} . Define $\mathbf{y} = \mathbf{x} - \mathbf{z}$.
2. Because both \mathbf{y} and \mathbf{w} are perpendicular to the hyperplane, we can rewrite $\mathbf{y} = v \frac{\mathbf{w}}{\|\mathbf{w}\|}$, where v is the Euclidean distance from \mathbf{z} to \mathbf{x} (what we are trying to get) and $\frac{\mathbf{w}}{\|\mathbf{w}\|}$ is the unit vector pointing at the direction of \mathbf{w} .

What is the distance from a sample \mathbf{z} to the hyperplane?



1. Let the point on the hyperplane closest to \mathbf{z} be \mathbf{x} . Define $\mathbf{y} = \mathbf{x} - \mathbf{z}$.
2. Because both \mathbf{y} and \mathbf{w} are perpendicular to the hyperplane, we can rewrite $\mathbf{y} = v \frac{\mathbf{w}}{\|\mathbf{w}\|}$, where v is the Euclidean distance from \mathbf{z} to \mathbf{x} (what we are trying to get) and $\frac{\mathbf{w}}{\|\mathbf{w}\|}$ is the unit vector pointing at the direction of \mathbf{w} .
3. Therefore, $\mathbf{z} = \mathbf{x} + v \frac{\mathbf{w}}{\|\mathbf{w}\|}$.

What is the distance from a sample \mathbf{z} to the hyperplane?




4. The prediction for \mathbf{z}

is then (substituting into linear classifier equation):

$$\begin{aligned} & \mathbf{w}^T \mathbf{z} + w_b \\ = & \mathbf{w}^T \left(\mathbf{x} + v \frac{\mathbf{w}}{\|\mathbf{w}\|} \right) + w_b \\ = & \mathbf{w}^T \mathbf{x} + v \frac{\mathbf{w}^T \mathbf{w}}{\|\mathbf{w}\|} + w_b = \underbrace{\mathbf{w}^T \mathbf{x} + w_b}_{=0, \text{ by definition}} + v \frac{\mathbf{w}^T \mathbf{w}}{\|\mathbf{w}\|} \\ = & v \frac{\mathbf{w}^T \mathbf{w}}{\|\mathbf{w}\|} = v \frac{\|\mathbf{w}\|^2}{\|\mathbf{w}\|} = v \|\mathbf{w}\|. \end{aligned}$$

1. Let the point on the hyperplane closest to \mathbf{z} be \mathbf{x} . Define $\mathbf{y} = \mathbf{x} - \mathbf{z}$.
2. Because both \mathbf{y} and \mathbf{w} are perpendicular to the hyperplane, we can rewrite $\mathbf{y} = v \frac{\mathbf{w}}{\|\mathbf{w}\|}$, where v is the Euclidean distance from \mathbf{z} to \mathbf{x} (what we are trying to get) and $\frac{\mathbf{w}}{\|\mathbf{w}\|}$ is the unit vector pointing at the direction of \mathbf{w} .
3. Therefore, $\mathbf{z} = \mathbf{x} + v \frac{\mathbf{w}}{\|\mathbf{w}\|}$.


What is the distance from a sample \mathbf{z} to the hyperplane?



1. Let the point on the hyperplane closest to \mathbf{z} be \mathbf{x} . Define $\mathbf{y} = \mathbf{x} - \mathbf{z}$.
2. Because both \mathbf{y} and \mathbf{w} are perpendicular to the hyperplane, we can rewrite $\mathbf{y} = v \frac{\mathbf{w}}{\|\mathbf{w}\|}$, where v is the Euclidean distance from \mathbf{z} to \mathbf{x} (what we are trying to get) and $\frac{\mathbf{w}}{\|\mathbf{w}\|}$ is the unit vector pointing at the direction of \mathbf{w} .
3. Therefore, $\mathbf{z} = \mathbf{x} + v \frac{\mathbf{w}}{\|\mathbf{w}\|}$.
4. The prediction for \mathbf{z} is then (substituting into linear classifier equation):
5. Finally, $v = \overbrace{\mathbf{w}^T \mathbf{z} + w_b}^{\text{prediction}} / \|\mathbf{w}\|$.

$$\begin{aligned} & \mathbf{w}^T \mathbf{z} + w_b \\ = & \mathbf{w}^T \left(\mathbf{x} + v \frac{\mathbf{w}}{\|\mathbf{w}\|} \right) + w_b \\ = & \mathbf{w}^T \mathbf{x} + v \frac{\mathbf{w}^T \mathbf{w}}{\|\mathbf{w}\|} + w_b = \underbrace{\mathbf{w}^T \mathbf{x} + w_b}_{=0, \text{ by definition}} + v \frac{\mathbf{w}^T \mathbf{w}}{\|\mathbf{w}\|} \\ = & v \frac{\mathbf{w}^T \mathbf{w}}{\|\mathbf{w}\|} = v \frac{\|\mathbf{w}\|^2}{\|\mathbf{w}\|} = v \|\mathbf{w}\|. \end{aligned}$$

What is the distance from a sample \mathbf{z} to the hyperplane?



4. The prediction for \mathbf{z}

is then (substituting into linear classifier equation):

$$\begin{aligned} & \mathbf{w}^T \mathbf{z} + w_b \\ = & \mathbf{w}^T \left(\mathbf{x} + v \frac{\mathbf{w}}{\|\mathbf{w}\|} \right) + w_b \\ = & \mathbf{w}^T \mathbf{x} + v \frac{\mathbf{w}^T \mathbf{w}}{\|\mathbf{w}\|} + w_b = \underbrace{\mathbf{w}^T \mathbf{x} + w_b}_{=0, \text{ by definition}} + v \frac{\mathbf{w}^T \mathbf{w}}{\|\mathbf{w}\|} \\ = & v \frac{\mathbf{w}^T \mathbf{w}}{\|\mathbf{w}\|} = v \frac{\|\mathbf{w}\|^2}{\|\mathbf{w}\|} = v \|\mathbf{w}\|. \end{aligned}$$

1. Let the point on the hyperplane closest to \mathbf{z} be \mathbf{x} .

Define $\mathbf{y} = \mathbf{x} - \mathbf{z}$.


2. Because both \mathbf{y} and \mathbf{w} are perpendicular to the hyperplane, we can rewrite $\mathbf{y} = v \frac{\mathbf{w}}{\|\mathbf{w}\|}$, where v is the Euclidean distance from \mathbf{z} to \mathbf{x} (what we are trying to get) and $\frac{\mathbf{w}}{\|\mathbf{w}\|}$ is the unit vector pointing at the direction of \mathbf{w} .

3. Therefore, $\mathbf{z} = \mathbf{x} + v \frac{\mathbf{w}}{\|\mathbf{w}\|}$.

5. Finally, $v = \overbrace{\mathbf{w}^T \mathbf{z} + w_b}^{\text{prediction}} / \|\mathbf{w}\|$.


6. **Conclusion:** a sample \mathbf{z} 's distance to a hyperplane $\mathbf{w}^T \mathbf{x} + w_b = 0$ is $d / \|\mathbf{w}\|$ **if and only if** the prediction for it $\mathbf{w}^T \mathbf{z} + w_b$ is $\pm d$. (The sign ahead of d depends on which side the sample is on.)

Hard margin linear SVM (for two linearly separable classes)




- All samples of Classes $+1$ and -1 are above and below the hyperplane, respectively.

Hard margin linear SVM (for two linearly separable classes)




- ▶ All samples of Classes $+1$ and -1 are above and below the hyperplane, respectively.
- ▶ For Class $+1$, denote the distance from the sample(s) closest to the hyperplane as $d_1/\|\mathbf{w}\|$ ($d_1 > 0$).

Hard margin linear SVM (for two linearly separable classes)




- ▶ All samples of Classes $+1$ and -1 are above and below the hyperplane, respectively.
- ▶ For Class $+1$, denote the distance from the sample(s) closest to the hyperplane as $d_1 / \|\mathbf{w}\|$ ($d_1 > 0$).
- ▶ Using the conclusion from previous slide, the prediction $\mathbf{w}^T \mathbf{x} + w_b$ for any sample \mathbf{x} of Class $+1$ is thus at least d_1 : $\mathbf{w}^T \mathbf{x} + w_b \geq d_1$.

Hard margin linear SVM (for two linearly separable classes)



- ▶ All samples of Classes $+1$ and -1 are above and below the hyperplane, respectively.
- ▶ For Class $+1$, denote the distance from the sample(s) closest to the hyperplane as $d_1 / \|\mathbf{w}\|$ ($d_1 > 0$).
- ▶ Using the conclusion from previous slide, the prediction $\mathbf{w}^T \mathbf{x} + w_b$ for any sample \mathbf{x} of Class $+1$ is thus at least d_1 : $\mathbf{w}^T \mathbf{x} + w_b \geq d_1$.
- ▶ Similarly, for Class -1 , we have $\mathbf{w}^T \mathbf{x} + w_b \leq -d_2$, where d_2 is the minimal distance. (Changes: $-$ and \leq)


Hard margin linear SVM (for two linearly separable classes)



- ▶ All samples of Classes $+1$ and -1 are above and below the hyperplane, respectively.
- ▶ For Class $+1$, denote the distance from the sample(s) closest to the hyperplane as $d_1 / \|\mathbf{w}\|$ ($d_1 > 0$).
- ▶ Using the conclusion from previous slide, the prediction $\mathbf{w}^T \mathbf{x} + w_b$ for any sample \mathbf{x} of Class $+1$ is thus at least d_1 : $\mathbf{w}^T \mathbf{x} + w_b \geq d_1$.
- ▶ Similarly, for Class -1 , we have $\mathbf{w}^T \mathbf{x} + w_b \leq -d_2$, where d_2 is the minimal distance. (Changes: $-$ and \leq)

- ▶ The idea of an SVM is to find a direction (defined by \mathbf{w}) along which closest samples of both classes are apart the most.


Hard margin linear SVM (for two linearly separable classes)



- ▶ All samples of Classes $+1$ and -1 are above and below the hyperplane, respectively.
- ▶ For Class $+1$, denote the distance from the sample(s) closest to the hyperplane as $d_1 / \|\mathbf{w}\|$ ($d_1 > 0$).
- ▶ Using the conclusion from previous slide, the prediction $\mathbf{w}^T \mathbf{x} + w_b$ for any sample \mathbf{x} of Class $+1$ is thus at least d_1 : $\mathbf{w}^T \mathbf{x} + w_b \geq d_1$.
- ▶ Similarly, for Class -1 , we have $\mathbf{w}^T \mathbf{x} + w_b \leq -d_2$, where d_2 is the minimal distance. (Changes: $-$ and \leq)

- ▶ The idea of an SVM is to find a direction (defined by \mathbf{w}) along which closest samples of both classes are apart the most.
- ▶ Hence, we want to maximize $\frac{d_1}{\|\mathbf{w}\|} + \frac{d_2}{\|\mathbf{w}\|}$, known as the **margin**.

Hard margin linear SVM (for two linearly separable classes)




- ▶ All samples of Classes $+1$ and -1 are above and below the hyperplane, respectively.
- ▶ For Class $+1$, denote the distance from the sample(s) closest to the hyperplane as $d_1 / \|\mathbf{w}\|$ ($d_1 > 0$).
- ▶ Using the conclusion from previous slide, the prediction $\mathbf{w}^T \mathbf{x} + w_b$ for any sample \mathbf{x} of Class $+1$ is thus at least d_1 : $\mathbf{w}^T \mathbf{x} + w_b \geq d_1$.
- ▶ Similarly, for Class -1 , we have $\mathbf{w}^T \mathbf{x} + w_b \leq -d_2$, where d_2 is the minimal distance. (Changes: $-$ and \leq)

- ▶ The idea of an SVM is to find a direction (defined by \mathbf{w}) along which closest samples of both classes are apart the most.
- ▶ Hence, we want to maximize $\frac{d_1}{\|\mathbf{w}\|} + \frac{d_2}{\|\mathbf{w}\|}$, known as the **margin**.

- ▶ Finally:


$$\begin{cases} \max & \frac{d_1}{\|\mathbf{w}\|} + \frac{d_2}{\|\mathbf{w}\|} \\ \text{s.t.} & \mathbf{w}^T \mathbf{x} + w_b - d_1 \geq 0, \forall \mathbf{x} \in C_{+1} \\ & \mathbf{w}^T \mathbf{x} + w_b + d_2 \geq 0, \forall \mathbf{x} \in C_{-1} \end{cases}$$

Hard margin linear SVM (cond.)




► We prefer $d_1 = d_2$: both classes are equal.

Hard margin linear SVM (cond.)



- ▶ We prefer $d_1 = d_2$: both classes are equal.
- ▶ Since d_1 and d_2 are constants, we can let them be 1.


Hard margin linear SVM (cond.)



- ▶ We prefer $d_1 = d_2$: both classes are equal.
- ▶ Since d_1 and d_2 are constants, we can let them be 1.
- ▶ Leveraging the label $y_k \in \{+1, -1\}$, we have a concise form:


$$\begin{cases} \max & \frac{2}{\|\mathbf{w}\|} \\ \text{s.t.} & y_k(\mathbf{w}^T \mathbf{x}_k + w_b) \geq 1, \forall \mathbf{x}_k \in C_{+1} \cup C_{-1}. \end{cases}$$

Hard margin linear SVM (cond.)



- ▶ We prefer $d_1 = d_2$: both classes are equal.
- ▶ Since d_1 and d_2 are constants, we can let them be 1.
- ▶ Leveraging the label $y_k \in \{+1, -1\}$, we have a concise form:
$$\begin{cases} \max & \frac{2}{\|\mathbf{w}\|} \\ \text{s.t.} & y_k(\mathbf{w}^T \mathbf{x}_k + w_b) \geq 1, \forall \mathbf{x}_k \in C_{+1} \cup C_{-1}. \end{cases}$$
- ▶ Maximizing $\frac{2}{\|\mathbf{w}\|}$ is equivalent to minimizing $\frac{\|\mathbf{w}\|}{2}$.

Hard margin linear SVM (cond.)




- ▶ We prefer $d_1 = d_2$: both classes are equal.
- ▶ Since d_1 and d_2 are constants, we can let them be 1.
- ▶ Leveraging the label $y_k \in \{+1, -1\}$, we have a concise form:

$$\begin{cases} \max & \frac{2}{\|\mathbf{w}\|} \\ \text{s.t.} & y_k(\mathbf{w}^T \mathbf{x}_k + w_b) \geq 1, \forall \mathbf{x}_k \in C_{+1} \cup C_{-1}. \end{cases}$$

- ▶ Maximizing $\frac{2}{\|\mathbf{w}\|}$ is equivalent to minimizing $\frac{\|\mathbf{w}\|}{2}$.
- ▶ Finally, we transform it into a quadratic programming problem (**the primal form of SVMs**):

$$\begin{cases} \min & \frac{1}{2} \|\mathbf{w}\|^2 = \frac{1}{2} \mathbf{w}^T \mathbf{w} \\ \text{s.t.} & y_k(\mathbf{w}^T \mathbf{x}_k + w_b) \geq 1, \forall \mathbf{x}_k. \end{cases}$$

Hard margin linear SVM (cond.)



- ▶ We prefer $d_1 = d_2$: both classes are equal.
- ▶ Since d_1 and d_2 are constants, we can let them be 1.
- ▶ Leveraging the label $y_k \in \{+1, -1\}$, we have a concise form:
$$\begin{cases} \max & \frac{2}{\|\mathbf{w}\|} \\ \text{s.t.} & y_k(\mathbf{w}^T \mathbf{x}_k + w_b) \geq 1, \forall \mathbf{x}_k \in C_{+1} \cup C_{-1}. \end{cases}$$
- ▶ Maximizing $\frac{2}{\|\mathbf{w}\|}$ is equivalent to minimizing $\frac{\|\mathbf{w}\|}{2}$.
- ▶ Finally, we transform it into a quadratic programming problem (**the primal form of SVMs**):
$$\begin{cases} \min & \frac{1}{2} \|\mathbf{w}\|^2 = \frac{1}{2} \mathbf{w}^T \mathbf{w} \\ \text{s.t.} & y_k(\mathbf{w}^T \mathbf{x}_k + w_b) \geq 1, \forall \mathbf{x}_k. \end{cases}$$
- ▶ Why square $\|\mathbf{w}\|$?

Recap: the Karush-Kuhn-Tucker (KKT) conditions

- Given a nonlinear optimization problem

$$\begin{cases} \min & f(\mathbf{x}) \\ s.t. & h_k(\mathbf{x}) \geq 0, \forall k \in [1..K], \end{cases}$$

where \mathbf{x} is a vector, and $h_k(\cdot)$ is linear, its Lagrange multiplier (or Lagrangian) is:

$$L(\mathbf{x}, \lambda) = f(\mathbf{x}) - \sum_{k=1}^K \lambda_k h_k(\mathbf{x})$$

Recap: the Karush-Kuhn-Tucker (KKT) conditions

- Given a nonlinear optimization problem

$$\begin{cases} \min & f(\mathbf{x}) \\ s.t. & h_k(\mathbf{x}) \geq 0, \forall k \in [1..K], \end{cases}$$

where \mathbf{x} is a vector, and $h_k(\cdot)$ is linear, its Lagrange multiplier (or Lagrangian) is:

$$L(\mathbf{x}, \lambda) = f(\mathbf{x}) - \sum_{k=1}^K \lambda_k h_k(\mathbf{x})$$

- The necessary conditions that the problem above has a solution are KKT conditions:

$$\begin{cases} \frac{\partial L}{\partial \mathbf{x}} = \mathbf{0}, \\ \lambda_k \geq 0, & \forall k \in [1..K] \\ \lambda_k h_k(\mathbf{x}) = 0, & \forall k \in [1..K] \end{cases}$$

Properties of hard margin linear SVM

The KKT condition to the SVM problem is

$$\begin{cases} A : \frac{\partial L}{\partial w} = \mathbf{0}, \\ B : \frac{\partial L}{\partial w_b} = 0, \\ C : \lambda_k \geq 0, & \forall k \in [1..K] \\ D : \lambda_k [y_k (\mathbf{w}^T \mathbf{x}_k + w_b) - 1] = 0, & \forall k \in [1..K] \end{cases}$$

Properties of hard margin linear SVM

The KKT condition to the SVM problem is

$$\begin{cases} A : \frac{\partial L}{\partial \mathbf{w}} = \mathbf{0}, \\ B : \frac{\partial L}{\partial w_b} = 0, \\ C : \lambda_k \geq 0, & \forall k \in [1..K] \\ D : \lambda_k [y_k (\mathbf{w}^T \mathbf{x}_k + w_b) - 1] = 0, & \forall k \in [1..K] \end{cases}$$

From Eqs. A and B,

$$\frac{\partial L}{\partial \mathbf{w}} = \mathbf{w} - \sum_{k=1}^K \lambda_k y_k \mathbf{x}_k \Rightarrow \mathbf{w} = \sum_{k=1}^K \lambda_k y_k \mathbf{x}_k$$
$$\frac{\partial L}{\partial w_b} = \sum_{k=1}^K \lambda_k y_k = 0$$

Properties of hard margin linear SVM

The KKT condition to the SVM problem is

$$\begin{cases} A : \frac{\partial L}{\partial \mathbf{w}} = \mathbf{0}, \\ B : \frac{\partial L}{\partial w_b} = 0, \\ C : \lambda_k \geq 0, & \forall k \in [1..K] \\ D : \lambda_k [y_k (\mathbf{w}^T \mathbf{x}_k + w_b) - 1] = 0, & \forall k \in [1..K] \end{cases}$$

From Eqs. A and B,


$$\begin{aligned} \frac{\partial L}{\partial \mathbf{w}} = \mathbf{w} - \sum_{k=1}^K \lambda_k y_k \mathbf{x}_k &\Rightarrow \mathbf{w} = \sum_{k=1}^K \lambda_k y_k \mathbf{x}_k \\ \frac{\partial L}{\partial w_b} = \sum_{k=1}^K \lambda_k y_k &= 0 \end{aligned}$$

Because λ_k is either positive or 0, the solution of the SVM problem is only associated with samples whose $\lambda_k \neq 0$. Denote them as $N_s = \{\mathbf{x}_k | \lambda_k \neq 0, k \in [1..K]\}$.

Properties of hard margin linear SVM (cont.)

- Therefore, Eq. A can be rewritten into

$$\mathbf{w} = \sum_{\mathbf{x}_k \in N_s} \lambda_k y_k \mathbf{x}_k$$




Properties of hard margin linear SVM (cont.)

- Therefore, Eq. A can be rewritten into

$$\mathbf{w} = \sum_{\mathbf{x}_k \in N_s} \lambda_k y_k \mathbf{x}_k$$

- The samples $\mathbf{x}_k \in N_s$ collectively determine the \mathbf{w} , and thus called **support vectors**, supporting the solution.




Properties of hard margin linear SVM (cont.)

- Therefore, Eq. A can be rewritten into

$$\mathbf{w} = \sum_{\mathbf{x}_k \in N_s} \lambda_k y_k \mathbf{x}_k$$

- The samples $\mathbf{x}_k \in N_s$ collectively determine the \mathbf{w} , and thus called **support vectors**, supporting the solution.
- The support vectors also have an interesting “visual” properties. From Eq. D, we have $\lambda_k [y_k (\mathbf{w}^T \mathbf{x}_k + w_b) - 1] = 0$. Because for $\mathbf{x}_k \in N_s$, $\lambda_k > 0$, then $y_k (\mathbf{w}^T \mathbf{x}_k + w_b) = 1$.




Properties of hard margin linear SVM (cont.)

- Therefore, Eq. A can be rewritten into

$$\mathbf{w} = \sum_{\mathbf{x}_k \in N_s} \lambda_k y_k \mathbf{x}_k$$

- The samples $\mathbf{x}_k \in N_s$ collectively determine the \mathbf{w} , and thus called **support vectors**, supporting the solution.
- The support vectors also have an interesting “visual” properties. From Eq. D, we have $\lambda_k [y_k (\mathbf{w}^T \mathbf{x}_k + w_b) - 1] = 0$. Because for $\mathbf{x}_k \in N_s$, $\lambda_k > 0$, then $y_k (\mathbf{w}^T \mathbf{x}_k + w_b) = 1$.
- Given that $y_k \in \{+1, -1\}$, we have $\mathbf{w}^T \mathbf{x}_k + w_b = \pm 1$. They support the **gutters**.



The dual form of an SVM

1. Given a nonlinear optimization problem in the **primal** form

$$\begin{cases} \min & f(\mathbf{x}) \\ s.t. & h_k(\mathbf{x}) \geq 0, \forall k \in [1..K], \end{cases}$$

The dual form of an SVM

1. Given a nonlinear optimization problem in the **primal** form

$$\begin{cases} \min & f(\mathbf{x}) \\ s.t. & h_k(\mathbf{x}) \geq 0, \forall k \in [1..K], \end{cases}$$

The dual form of an SVM

1. Given a nonlinear optimization problem in the **primal** form

$$\begin{cases} \min & f(\mathbf{x}) \\ s.t. & h_k(\mathbf{x}) \geq 0, \forall k \in [1..K], \end{cases}$$

2. its **dual form** is

$$\begin{cases} \max & L(\mathbf{x}, \lambda) = f(\mathbf{x}) - \sum_{k=1}^K \lambda_k h_k(\mathbf{x}) \\ s.t. & \lambda_k \geq 0, \forall k \in [1..K], \\ & \nabla L = \mathbf{0} \end{cases}$$

The dual form of an SVM

1. Given a nonlinear optimization problem in the **primal** form

$$\begin{cases} \min & f(\mathbf{x}) \\ s.t. & h_k(\mathbf{x}) \geq 0, \forall k \in [1..K], \end{cases}$$

2. its **dual form** is

$$\begin{cases} \max & L(\mathbf{x}, \lambda) = f(\mathbf{x}) - \sum_{k=1}^K \lambda_k h_k(\mathbf{x}) \\ s.t. & \lambda_k \geq 0, \forall k \in [1..K], \\ & \nabla L = \mathbf{0} \end{cases}$$

The dual form of an SVM

1. Given a nonlinear optimization problem in the **primal** form
3. Thus for a **primal** SVM problem

$$\begin{cases} \min & f(\mathbf{x}) \\ s.t. & h_k(\mathbf{x}) \geq 0, \forall k \in [1..K], \end{cases}$$

$$\begin{cases} \min & \frac{1}{2} \|\mathbf{w}\|^2 = \frac{1}{2} \mathbf{w}^T \mathbf{w} \\ s.t. & y_k(\mathbf{w}^T \mathbf{x}_k + w_b) \geq 1, \forall \mathbf{x}_k. \end{cases}$$

2. its **dual form** is

$$\begin{cases} \max & L(\mathbf{x}, \lambda) = f(\mathbf{x}) - \sum_{k=1}^K \lambda_k h_k(\mathbf{x}) \\ s.t. & \lambda_k \geq 0, \forall k \in [1..K], \\ & \nabla L = \mathbf{0} \end{cases}$$

The dual form of an SVM

1. Given a nonlinear optimization problem in the **primal** form
3. Thus for a **primal** SVM problem

$$\begin{cases} \min & f(\mathbf{x}) \\ s.t. & h_k(\mathbf{x}) \geq 0, \forall k \in [1..K], \end{cases}$$

$$\begin{cases} \min & \frac{1}{2} \|\mathbf{w}\|^2 = \frac{1}{2} \mathbf{w}^T \mathbf{w} \\ s.t. & y_k(\mathbf{w}^T \mathbf{x}_k + w_b) \geq 1, \forall \mathbf{x}_k. \end{cases}$$

2. its **dual form** is

$$\begin{cases} \max & L(\mathbf{x}, \lambda) = f(\mathbf{x}) - \sum_{k=1}^K \lambda_k h_k(\mathbf{x}) \\ s.t. & \lambda_k \geq 0, \forall k \in [1..K], \\ & \nabla L = \mathbf{0} \end{cases}$$

The dual form of an SVM

1. Given a nonlinear optimization problem in the **primal** form

$$\begin{cases} \min & f(\mathbf{x}) \\ \text{s.t.} & h_k(\mathbf{x}) \geq 0, \forall k \in [1..K], \end{cases}$$

2. its **dual form** is

$$\begin{cases} \max & L(\mathbf{x}, \lambda) = f(\mathbf{x}) - \sum_{k=1}^K \lambda_k h_k(\mathbf{x}) \\ \text{s.t.} & \lambda_k \geq 0, \forall k \in [1..K], \\ & \nabla L = \mathbf{0} \end{cases}$$

3. Thus for a **primal** SVM problem

$$\begin{cases} \min & \frac{1}{2} \|\mathbf{w}\|^2 = \frac{1}{2} \mathbf{w}^T \mathbf{w} \\ \text{s.t.} & y_k(\mathbf{w}^T \mathbf{x}_k + w_b) \geq 1, \forall \mathbf{x}_k. \end{cases}$$

4. its **dual form** is

$$\begin{cases} \max & \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{k=1}^K \lambda_k (y_k(\mathbf{w}^T \mathbf{x}_k + w_b) - 1) \\ \text{s.t.} & \lambda_k \geq 0, \forall k \in [1..K], \\ & \mathbf{w} = \sum_{k=1}^K \lambda_k y_k \mathbf{x}_k \quad (\text{from } \frac{\partial L}{\partial \mathbf{w}} = 0), \\ & \sum_{k=1}^K \lambda_k y_k = 0 \quad (\text{from } \frac{\partial L}{\partial w_b} = 0) \end{cases}$$

The dual form of an SVM (cond.)

$$\left\{ \begin{array}{l} \max \quad \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{k=1}^K \lambda_k (y_k (\mathbf{w}^T \mathbf{x}_k + w_b) - 1) \\ \text{s.t.} \quad \lambda_k \geq 0, \forall k \in [1..K], \\ \mathbf{w} = \sum_{k=1}^K \lambda_k y_k \mathbf{x}_k \quad (\text{from } \frac{\partial L}{\partial \mathbf{w}} = 0), \\ \sum_{k=1}^K \lambda_k y_k = 0 \quad (\text{from } \frac{\partial L}{\partial w_b} = 0) \end{array} \right.$$

Substituting \mathbf{w} with $\sum_{k=1}^K \lambda_k y_k \mathbf{x}_k$, the objective function becomes:

$$L = -\frac{1}{2} \sum_{i=1}^K \sum_{j=1}^K \lambda_i \lambda_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j + \sum_{k=1}^K \lambda_k$$

Thus, the new **dual form** is:

$$\left\{ \begin{array}{l} \max \quad -\frac{1}{2} \sum_{i=1}^K \sum_{j=1}^K \lambda_i \lambda_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j + \sum_{k=1}^K \lambda_k \\ \text{s.t.} \quad \lambda_k \geq 0, \forall k \in [1..K], \\ \sum_{k=1}^K \lambda_k y_k = 0 \end{array} \right.$$

- The number of unknowns to solve drops from n features to K samples.

The dual form of an SVM (cond.)

$$\left\{ \begin{array}{l} \max \quad \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{k=1}^K \lambda_k (y_k (\mathbf{w}^T \mathbf{x}_k + w_b) - 1) \\ \text{s.t.} \quad \lambda_k \geq 0, \forall k \in [1..K], \\ \mathbf{w} = \sum_{k=1}^K \lambda_k y_k \mathbf{x}_k \quad (\text{from } \frac{\partial L}{\partial \mathbf{w}} = 0), \\ \sum_{k=1}^K \lambda_k y_k = 0 \quad (\text{from } \frac{\partial L}{\partial w_b} = 0) \end{array} \right.$$

Substituting \mathbf{w} with $\sum_{k=1}^K \lambda_k y_k \mathbf{x}_k$, the objective function becomes:

$$L = -\frac{1}{2} \sum_{i=1}^K \sum_{j=1}^K \lambda_i \lambda_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j + \sum_{k=1}^K \lambda_k$$

Thus, the new **dual form** is:

$$\left\{ \begin{array}{l} \max \quad -\frac{1}{2} \sum_{i=1}^K \sum_{j=1}^K \lambda_i \lambda_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j + \sum_{k=1}^K \lambda_k \\ \text{s.t.} \quad \lambda_k \geq 0, \forall k \in [1..K], \\ \sum_{k=1}^K \lambda_k y_k = 0 \end{array} \right.$$

- ▶ The number of unknowns to solve drops from n features to K samples.
- ▶ Instead of finding \mathbf{w} , find K λ_k 's. (Is an SVM really non-parametric?)

The dual form of an SVM (cond.)

$$\left\{ \begin{array}{l} \max \quad \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{k=1}^K \lambda_k (y_k (\mathbf{w}^T \mathbf{x}_k + w_b) - 1) \\ \text{s.t.} \quad \lambda_k \geq 0, \forall k \in [1..K], \\ \mathbf{w} = \sum_{k=1}^K \lambda_k y_k \mathbf{x}_k \quad (\text{from } \frac{\partial L}{\partial \mathbf{w}} = 0), \\ \sum_{k=1}^K \lambda_k y_k = 0 \quad (\text{from } \frac{\partial L}{\partial w_b} = 0) \end{array} \right.$$

Substituting \mathbf{w} with $\sum_{k=1}^K \lambda_k y_k \mathbf{x}_k$, the objective function becomes:

$$L = -\frac{1}{2} \sum_{i=1}^K \sum_{j=1}^K \lambda_i \lambda_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j + \sum_{k=1}^K \lambda_k$$

Thus, the new **dual form** is:

$$\left\{ \begin{array}{l} \max \quad -\frac{1}{2} \sum_{i=1}^K \sum_{j=1}^K \lambda_i \lambda_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j + \sum_{k=1}^K \lambda_k \\ \text{s.t.} \quad \lambda_k \geq 0, \forall k \in [1..K], \\ \sum_{k=1}^K \lambda_k y_k = 0 \end{array} \right.$$

- ▶ The number of unknowns to solve drops from n features to K samples.
- ▶ Instead of finding \mathbf{w} , find K λ_k 's. (Is an SVM really non-parametric?)
- ▶ The new SVM: $g(\mathbf{x}) = \mathbf{w}^T \mathbf{x} = \sum_{k=1}^K \lambda_k y_k (\mathbf{x}^T \mathbf{x}_k) + w_b$.

The dual form of an SVM (cond.)

$$\left\{ \begin{array}{l} \max \quad \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{k=1}^K \lambda_k (y_k (\mathbf{w}^T \mathbf{x}_k + w_b) - 1) \\ \text{s.t.} \quad \lambda_k \geq 0, \forall k \in [1..K], \\ \mathbf{w} = \sum_{k=1}^K \lambda_k y_k \mathbf{x}_k \quad (\text{from } \frac{\partial L}{\partial \mathbf{w}} = 0), \\ \sum_{k=1}^K \lambda_k y_k = 0 \quad (\text{from } \frac{\partial L}{\partial w_b} = 0) \end{array} \right.$$

Substituting \mathbf{w} with $\sum_{k=1}^K \lambda_k y_k \mathbf{x}_k$, the objective function becomes:

$$L = -\frac{1}{2} \sum_{i=1}^K \sum_{j=1}^K \lambda_i \lambda_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j + \sum_{k=1}^K \lambda_k$$

Thus, the new **dual form** is:

$$\left\{ \begin{array}{l} \max \quad -\frac{1}{2} \sum_{i=1}^K \sum_{j=1}^K \lambda_i \lambda_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j + \sum_{k=1}^K \lambda_k \\ \text{s.t.} \quad \lambda_k \geq 0, \forall k \in [1..K], \\ \sum_{k=1}^K \lambda_k y_k = 0 \end{array} \right.$$

- ▶ The number of unknowns to solve drops from n features to K samples.
- ▶ Instead of finding \mathbf{w} , find K λ_k 's. (Is an SVM really non-parametric?)
- ▶ The new SVM: $g(\mathbf{x}) = \mathbf{w}^T \mathbf{x} = \sum_{k=1}^K \lambda_k y_k (\mathbf{x}^T \mathbf{x}_k) + w_b$.
- ▶ To store an SVM model, just store the support vectors \mathbf{x}_i 's, their labels y_i 's and weights λ_i 's, and the bias w_b .

Kernel tricks: achieving non-linearity on SVMs

- ▶ In the previous slides, any two samples “interact” with each other thru dot product, e.g., $\mathbf{x}_i^T \mathbf{x}_j$ (in training, between two samples) or $\mathbf{x}^T \mathbf{x}_k$ (in prediction, between a sample to be predicted and a support vector).

Kernel tricks: achieving non-linearity on SVMs

- ▶ In the previous slides, any two samples “interact” with each other thru dot product, e.g., $\mathbf{x}_i^T \mathbf{x}_j$ (in training, between two samples) or $\mathbf{x}^T \mathbf{x}_k$ (in prediction, between a sample to be predicted and a support vector).
- ▶ It can be expanded to any function, denoted as $\mathcal{K}(\mathbf{x}, \mathbf{y})$ (\mathbf{x} and \mathbf{y} are any two vectors of same dimension. Not the input and output of an estimator), between two vectors, known as the **kernel function** or **kernel tricks**.

Kernel tricks: achieving non-linearity on SVMs

- ▶ In the previous slides, any two samples “interact” with each other thru dot product, e.g., $\mathbf{x}_i^T \mathbf{x}_j$ (in training, between two samples) or $\mathbf{x}^T \mathbf{x}_k$ (in prediction, between a sample to be predicted and a support vector).
- ▶ It can be expanded to any function, denoted as $\mathcal{K}(\mathbf{x}, \mathbf{y})$ (\mathbf{x} and \mathbf{y} are any two vectors of same dimension. Not the input and output of an estimator), between two vectors, known as the **kernel function** or **kernel tricks**.
- ▶ Linear kernels: what we have seen so far in SVMs.

Kernel tricks: achieving non-linearity on SVMs

- ▶ In the previous slides, any two samples “interact” with each other thru dot product, e.g., $\mathbf{x}_i^T \mathbf{x}_j$ (in training, between two samples) or $\mathbf{x}^T \mathbf{x}_k$ (in prediction, between a sample to be predicted and a support vector).
- ▶ It can be expanded to any function, denoted as $\mathcal{K}(\mathbf{x}, \mathbf{y})$ (\mathbf{x} and \mathbf{y} are any two vectors of same dimension. Not the input and output of an estimator), between two vectors, known as the **kernel function** or **kernel tricks**.
- ▶ Linear kernels: what we have seen so far in SVMs.
- ▶ Polynomial kernels: $\mathcal{K}(\mathbf{x}, \mathbf{y}) = (\mathbf{x} \cdot \mathbf{y} + b)^p$ where $p \in \mathbb{Z}^+$ and $b \in \mathbb{R}$.


Kernel tricks: achieving non-linearity on SVMs

- ▶ In the previous slides, any two samples “interact” with each other thru dot product, e.g., $\mathbf{x}_i^T \mathbf{x}_j$ (in training, between two samples) or $\mathbf{x}^T \mathbf{x}_k$ (in prediction, between a sample to be predicted and a support vector).
- ▶ It can be expanded to any function, denoted as $\mathcal{K}(\mathbf{x}, \mathbf{y})$ (\mathbf{x} and \mathbf{y} are any two vectors of same dimension. Not the input and output of an estimator), between two vectors, known as the **kernel function** or **kernel tricks**.
- ▶ Linear kernels: what we have seen so far in SVMs.
- ▶ Polynomial kernels: $\mathcal{K}(\mathbf{x}, \mathbf{y}) = (\mathbf{x} \cdot \mathbf{y} + b)^p$ where $p \in \mathbb{Z}^+$ and $b \in \mathbb{R}$.
- ▶ Gaussian (radial basis function, RBF) kernels (that build contours around support vectors when \mathbf{y} is a support vector): $\mathcal{K}(\mathbf{x}, \mathbf{y}) = \exp(-\|\mathbf{x} - \mathbf{y}\|^2 / \sigma)$

Kernel tricks: achieving non-linearity on SVMs

- ▶ In the previous slides, any two samples “interact” with each other thru dot product, e.g., $\mathbf{x}_i^T \mathbf{x}_j$ (in training, between two samples) or $\mathbf{x}^T \mathbf{x}_k$ (in prediction, between a sample to be predicted and a support vector).
- ▶ It can be expanded to any function, denoted as $\mathcal{K}(\mathbf{x}, \mathbf{y})$ (\mathbf{x} and \mathbf{y} are any two vectors of same dimension. Not the input and output of an estimator), between two vectors, known as the **kernel function** or **kernel tricks**.
- ▶ Linear kernels: what we have seen so far in SVMs.
- ▶ Polynomial kernels: $\mathcal{K}(\mathbf{x}, \mathbf{y}) = (\mathbf{x} \cdot \mathbf{y} + b)^p$ where $p \in \mathbb{Z}^+$ and $b \in \mathbb{R}$.
- ▶ Gaussian (radial basis function, RBF) kernels (that build contours around support vectors when \mathbf{y} is a support vector): $\mathcal{K}(\mathbf{x}, \mathbf{y}) = \exp(-\|\mathbf{x} - \mathbf{y}\|^2 / \sigma)$
- ▶ Usually linear and Gaussian are good enough. A Gaussian kernel can be decomposed into many polynomial terms.

Transforming a nonlinearly separable problem to a linearly separable one



Source: Wikipedia/SVM.

Generalized Linear Classifier

- ▶ Let $f_1(\cdot), f_2(\cdot), \dots, f_P(\cdot)$ be P nonlinear functions where $f_p : \mathbb{R}^n \mapsto \mathbb{R}, \forall p \in [1..P]$.

Generalized Linear Classifier

- ▶ Let $f_1(\cdot), f_2(\cdot), \dots, f_P(\cdot)$ be P nonlinear functions where $f_p : \mathbb{R}^n \mapsto \mathbb{R}, \forall p \in [1..P]$.
- ▶ Then we can define a mapping from a feature vector $\mathbf{x} \in \mathbb{R}^n$ (the **input space**) to a vector in another space $\mathbf{z} = [f_1(\mathbf{x}), f_2(\mathbf{x}), \dots, f_P(\mathbf{x})]^T \in \mathbb{R}^P$, which is called the **feature space**.

Generalized Linear Classifier

- ▶ Let $f_1(\cdot), f_2(\cdot), \dots, f_P(\cdot)$ be P nonlinear functions where $f_p : \mathbb{R}^n \mapsto \mathbb{R}, \forall p \in [1..P]$.
- ▶ Then we can define a mapping from a feature vector $\mathbf{x} \in \mathbb{R}^n$ (the **input space**) to a vector in another space $\mathbf{z} = [f_1(\mathbf{x}), f_2(\mathbf{x}), \dots, f_P(\mathbf{x})]^T \in \mathbb{R}^P$, which is called the **feature space**.
- ▶ The problem then becomes finding the value P and the functions $f_p(\cdot)$ such that the two classes are linearly separable.

Generalized Linear Classifier

- ▶ Let $f_1(\cdot), f_2(\cdot), \dots, f_P(\cdot)$ be P nonlinear functions where $f_p : \mathbb{R}^n \mapsto \mathbb{R}, \forall p \in [1..P]$.
- ▶ Then we can define a mapping from a feature vector $\mathbf{x} \in \mathbb{R}^n$ (the **input space**) to a vector in another space $\mathbf{z} = [f_1(\mathbf{x}), f_2(\mathbf{x}), \dots, f_P(\mathbf{x})]^T \in \mathbb{R}^P$, which is called the **feature space**.
- ▶ The problem then becomes finding the value P and the functions $f_p(\cdot)$ such that the two classes are linearly separable.
- ▶ Once the space transform is done, we wanna find a weight vector $\mathbf{w} \in \mathbb{R}^P$ such that

$$\begin{cases} \mathbf{w}^T \mathbf{z} + w_b > 0 & \text{if } \mathbf{z} \in C_1 \\ \mathbf{w}^T \mathbf{z} + w_b < 0 & \text{if } \mathbf{z} \in C_2. \end{cases}$$

Generalized Linear Classifier

- ▶ Let $f_1(\cdot), f_2(\cdot), \dots, f_P(\cdot)$ be P nonlinear functions where $f_p : \mathbb{R}^n \mapsto \mathbb{R}, \forall p \in [1..P]$.
- ▶ Then we can define a mapping from a feature vector $\mathbf{x} \in \mathbb{R}^n$ (the **input space**) to a vector in another space $\mathbf{z} = [f_1(\mathbf{x}), f_2(\mathbf{x}), \dots, f_P(\mathbf{x})]^T \in \mathbb{R}^P$, which is called the **feature space**.
- ▶ The problem then becomes finding the value P and the functions $f_p(\cdot)$ such that the two classes are linearly separable.
- ▶ Once the space transform is done, we wanna find a weight vector $\mathbf{w} \in \mathbb{R}^P$ such that

$$\begin{cases} \mathbf{w}^T \mathbf{z} + w_b > 0 & \text{if } \mathbf{z} \in C_1 \\ \mathbf{w}^T \mathbf{z} + w_b < 0 & \text{if } \mathbf{z} \in C_2. \end{cases}$$


- ▶ Essentially, we are building a new hyperplane $g(\mathbf{x}) = 0$ such that $g(\mathbf{x}) = w_b + \sum_{p=1}^P w_p f_p(\mathbf{x})$. Instead of computing the weighted sum of elements of feature vector, we compute that of elements of the transformed vector.

Creating features from input features

- For example, $g(\mathbf{x}) = w_b + w_1x_1 + w_2x_2 + w_{12}x_1x_2 + w_{11}x_1^2 + w_{22}x_2^2$

$$\Phi : R^2 \rightarrow R^3$$

$$(x_1, x_2) \mapsto (z_1, z_2, z_3) := (x_1^2, \sqrt{2}x_1x_2, x_2^2)$$




Creating features from input features

- ▶ For example, $g(\mathbf{x}) = w_b + w_1x_1 + w_2x_2 + w_{12}x_1x_2 + w_{11}x_1^2 + w_{22}x_2^2$
- ▶ Here is another example,

$$\Phi : R^2 \rightarrow R^3$$

$$(x_1, x_2) \mapsto (z_1, z_2, z_3) := (x_1^2, \sqrt{2}x_1x_2, x_2^2)$$




Creating features from input features

- ▶ For example, $g(\mathbf{x}) = w_b + w_1x_1 + w_2x_2 + w_{12}x_1x_2 + w_{11}x_1^2 + w_{22}x_2^2$
- ▶ Here is another example,


$$\Phi : \mathbb{R}^2 \rightarrow \mathbb{R}^3$$

$$(x_1, x_2) \mapsto (z_1, z_2, z_3) := (x_1^2, \sqrt{2}x_1x_2, x_2^2)$$




- ▶ A good explanation on StackOverflow:
<https://stats.stackexchange.com/questions/46425/what-is-feature-space>

Soft margin linear SVM



- We could allow some samples to fall into the margin in exchange for wider margin on the remaining samples.

Soft margin linear SVM




- ▶ We could allow some samples to fall into the margin in exchange for wider margin on the remaining samples.
- ▶ Therefore, we have a new optimization problem:

$$\begin{cases} \min & \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{k=1}^K \xi_k \\ s.t. & y_k (\mathbf{w}^T \mathbf{x}_k + w_b) \geq 1 - \xi_k, \forall \mathbf{x}_k \\ & \xi_k \geq 0. \end{cases}$$

where C is a constant, and ξ_k is called a **slack** variable defined as $\max(0, 1 - y_i(\mathbf{w}^T \mathbf{x}_k + w_b))$.

Soft margin linear SVM




- ▶ We could allow some samples to fall into the margin in exchange for wider margin on the remaining samples.
- ▶ Therefore, we have a new optimization problem:

$$\begin{cases} \min & \frac{1}{2} \|w\|^2 + C \sum_{k=1}^K \xi_k \\ s.t. & y_k (w^T x_k + w_b) \geq 1 - \xi_k, \forall x_k \\ & \xi_k \geq 0. \end{cases}$$

where C is a constant, and ξ_k is called a **slack** variable defined as $\max(0, 1 - y_i (w^T x_k + w_b))$.

- ▶ Such SVM is called *soft-margin*.

Soft margin linear SVM




- ▶ We could allow some samples to fall into the margin in exchange for wider margin on the remaining samples.
- ▶ Therefore, we have a new optimization problem:

$$\begin{cases} \min & \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{k=1}^K \xi_k \\ \text{s.t.} & y_k (\mathbf{w}^T \mathbf{x}_k + w_b) \geq 1 - \xi_k, \forall \mathbf{x}_k \\ & \xi_k \geq 0. \end{cases}$$

where C is a constant, and ξ_k is called a **slack** variable defined as $\max(0, 1 - y_i(\mathbf{w}^T \mathbf{x}_k + w_b))$.

- ▶ Such SVM is called *soft-margin*.
- ▶ The constant C provides a balance between maximizing the margin and minimizing the quality, instead of quantity, of misclassification.

Soft margin linear SVM



- ▶ We could allow some samples to fall into the margin in exchange for wider margin on the remaining samples.
- ▶ Therefore, we have a new optimization problem:

$$\begin{cases} \min & \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{k=1}^K \xi_k \\ \text{s.t.} & y_k (\mathbf{w}^T \mathbf{x}_k + w_b) \geq 1 - \xi_k, \forall \mathbf{x}_k \\ & \xi_k \geq 0. \end{cases}$$

where C is a constant, and ξ_k is called a **slack** variable defined as $\max(0, 1 - y_i(\mathbf{w}^T \mathbf{x}_k + w_b))$.

- ▶ Such SVM is called *soft-margin*.
- ▶ The constant C provides a balance between maximizing the margin and minimizing the quality, instead of quantity, of misclassification.
- ▶ Next: How to find C and why is slack variable defined so.

Grid search for hyperparameters

- ▶ Hyperparameters: Parameters of a model that is not updated in training but set based on experience or arbitrarily.

Grid search for hyperparameters

- ▶ Hyperparameters: Parameters of a model that is not updated in training but set based on experience or arbitrarily.
- ▶ Grid search: Create a sequence of values for each hyperparameter and form a grid from them using Cartesian product. Then for each point on the grid, evaluate the performance of the model. Finally, use the one that yields the best performance.

Grid search for hyperparameters

- ▶ Hyperparameters: Parameters of a model that is not updated in training but set based on experience or arbitrarily.
- ▶ Grid search: Create a sequence of values for each hyperparameter and form a grid from them using Cartesian product. Then for each point on the grid, evaluate the performance of the model. Finally, use the one that yields the best performance.
- ▶ How to evaluate the performance of a classifier?

Test set

- ▶ It would be unfair to evaluate the performance of a classifier using samples seen by the model during training.

Test set

- ▶ It would be unfair to evaluate the performance of a classifier using samples seen by the model during training.
- ▶ Samples unseen in training and used to evaluate the performance of a model form the **test set**.

Test set

- ▶ It would be unfair to evaluate the performance of a classifier using samples seen by the model during training.
- ▶ Samples unseen in training and used to evaluate the performance of a model form the **test set**.
- ▶ So, from all your data, you split them into two groups **training set** and test set.

Test set

- ▶ It would be unfair to evaluate the performance of a classifier using samples seen by the model during training.
- ▶ Samples unseen in training and used to evaluate the performance of a model form the **test set**.
- ▶ So, from all your data, you split them into two groups **training set** and test set.
- ▶ But, is just one test set good?

Cross-validation

- ▶ Cross validation (CV): split your data into many pairs of training and test sets. Then evaluate the performance of the classifier on each pair. Usually the test sets do not overlap. And, of course, the training and test sets in each pair do not overlap.

Cross-validation

- ▶ Cross validation (CV): split your data into many pairs of training and test sets. Then evaluate the performance of the classifier on each pair. Usually the test sets do not overlap. And, of course, the training and test sets in each pair do not overlap.
- ▶ k-fold CV: Split all data into k folds, equal-size and **non-overlapping**. In each round the CV, use $k - 1$ folds for training and the rest one fold for test. Then rotate on the test set. Stop after every fold has been used as test set exactly k times.

Cross-validation

- ▶ Cross validation (CV): split your data into many pairs of training and test sets. Then evaluate the performance of the classifier on each pair. Usually the test sets do not overlap. And, of course, the training and test sets in each pair do not overlap.
- ▶ k-fold CV: Split all data into k folds, equal-size and **non-overlapping**. In each round the CV, use $k - 1$ folds for training and the rest one fold for test. Then rotate on the test set. Stop after every fold has been used as test set exactly k times.
- ▶ leave-N-out CV (LNOCV): A special case of k-fold CV that only N samples are the test set. When $N = 1$, it becomes leave-one-out CV (LOOCV).

The slack variable and hinge loss

- ▶ What is the $\xi = \max(0, 1 - y_i(\mathbf{w}^T \mathbf{x} + w_b))$ when a sample \mathbf{x} is correctly classified?

The slack variable and hinge loss

- ▶ What is the $\xi = \max(0, 1 - y_i(\mathbf{w}^T \mathbf{x} + w_b))$ when a sample \mathbf{x} is correctly classified?
- ▶ It's zero.

The slack variable and hinge loss

- ▶ What is the $\xi = \max(0, 1 - y_i(\mathbf{w}^T \mathbf{x} + w_b))$ when a sample \mathbf{x} is correctly classified?
- ▶ It's zero.
- ▶ In that case, the constraint is the same as that for hard margin linear SVMs:
 $y_k(\mathbf{w}^T \mathbf{x} + w_b) \geq 0.$

The slack variable and hinge loss

- ▶ What is the $\xi = \max(0, 1 - y_i(\mathbf{w}^T \mathbf{x} + w_b))$ when a sample \mathbf{x} is correctly classified?
- ▶ It's zero.
- ▶ In that case, the constraint is the same as that for hard margin linear SVMs:
 $y_k(\mathbf{w}^T \mathbf{x} + w_b) \geq 0.$
- ▶ The expression $\max(0, 1 - y \cdot \hat{y})$ where $y \in \{+1, -1\}$ is the ground truth label and \hat{y} is prediction for a classifier, is called a **hinge loss**. It's “hinge” because as long as the classification is correct, the loss/error is (capped at) 0.